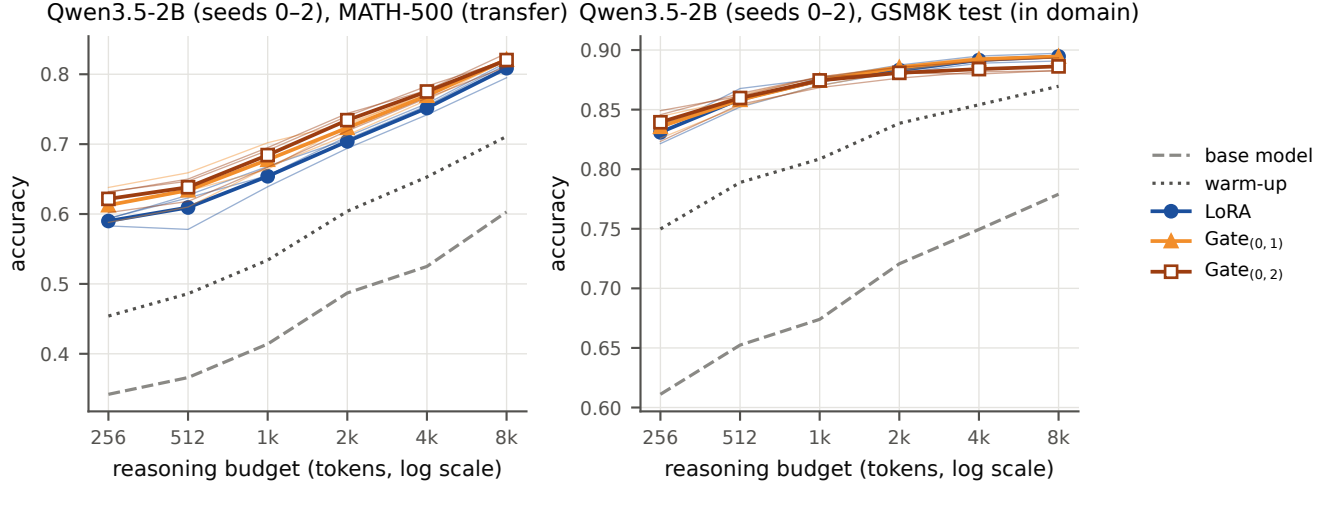
Qwen3.5-2B (seeds 0–2), MATH-500 (transfer)
Qwen3.5-2B (seeds 0–2), GSM8K test (in domain)
accuracy
reasoning budget (tokens, log scale)
256
512
1k
2k
4k
8k
0.4
0.5
0.6
0.7
0.8
0.60
0.65
0.70
0.75
0.80
0.85
0.90
base model
warm-up
LoRA
Gate(0, 1)
Gate(0, 2)