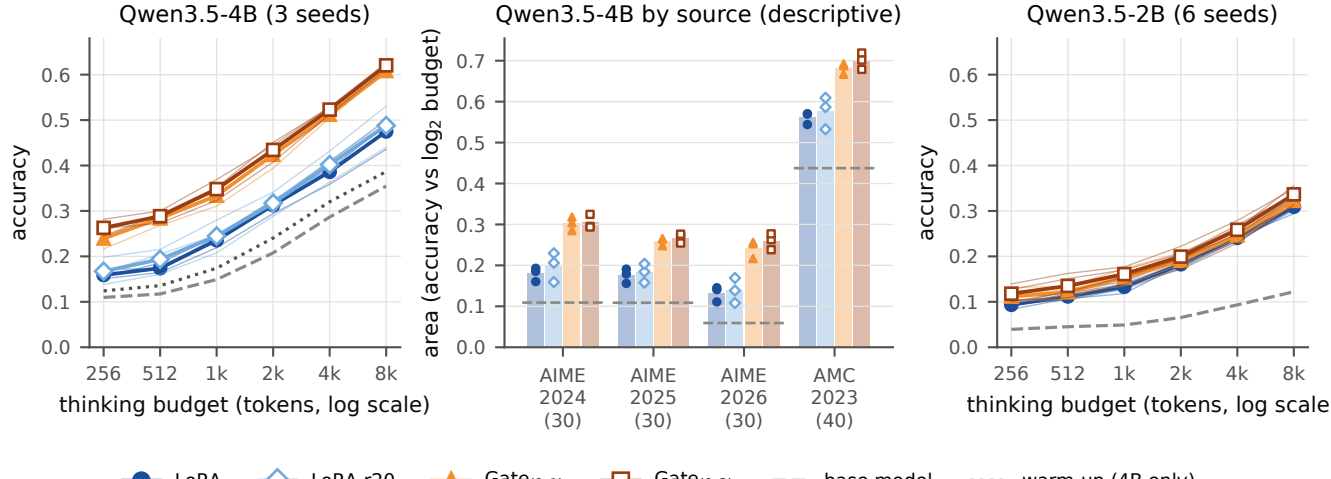
Qwen3.5-4B (3 seeds)
accuracy
0.0
0.1
0.2
0.3
0.4
0.5
0.6
256
512
1k
2k
4k
8k
thinking budget (tokens, log scale)
Qwen3.5-4B by source (descriptive)
area (accuracy vs log2 budget)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
AIME 2024 (30)
AIME 2025 (30)
AIME 2026 (30)
AMC 2023 (40)
Qwen3.5-2B (6 seeds)
accuracy
0.0
0.1
0.2
0.3
0.4
0.5
0.6
256
512
1k
2k
4k
8k
thinking budget (tokens, log scale)
LoRA
LoRA-r20
Gate(0, 1)
Gate(0, 2)
base model
warm-up (4B only)