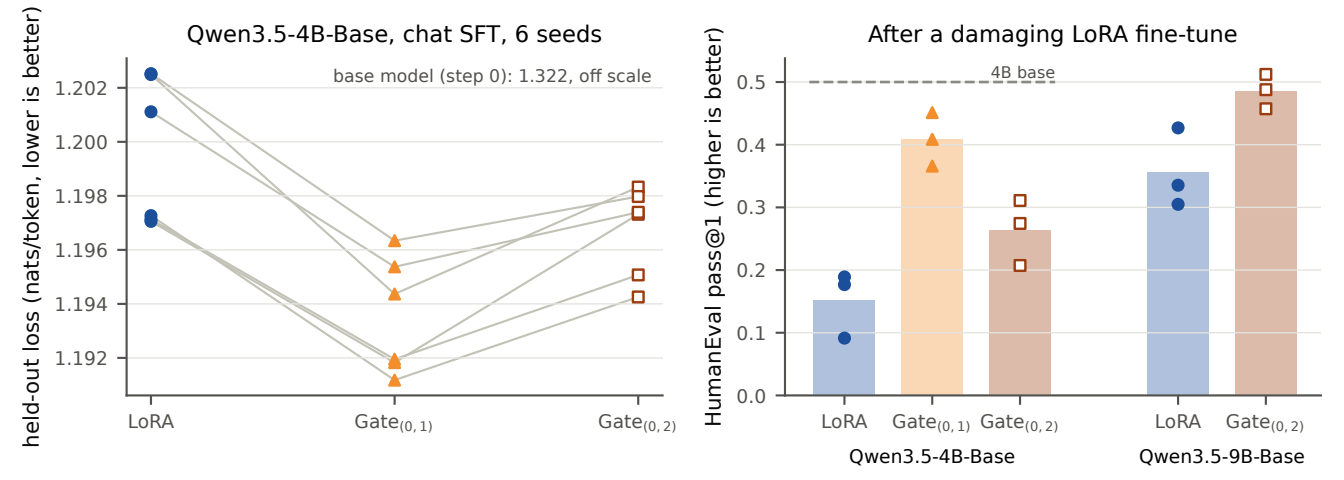
Qwen3.5-4B-Base, chat SFT, 6 seeds
base model (step 0): 1.322, off scale
held-out loss (nats/token, lower is better)
1.202
1.200
1.198
1.196
1.194
1.192
LoRA
Gate(0, 1)
Gate(0, 2)
After a damaging LoRA fine-tune
4B base
HumanEval pass@1 (higher is better)
0.5
0.4
0.3
0.2
0.1
0.0
LoRA
Gate(0, 1)
Gate(0, 2)
LoRA
Gate(0, 2)
Qwen3.5-4B-Base
Qwen3.5-9B-Base